



## Agenda


- Recap of delivery properties
- Implementing causal broadcast
- Use of causality
- Chandy-Lamport snapshot algorithm



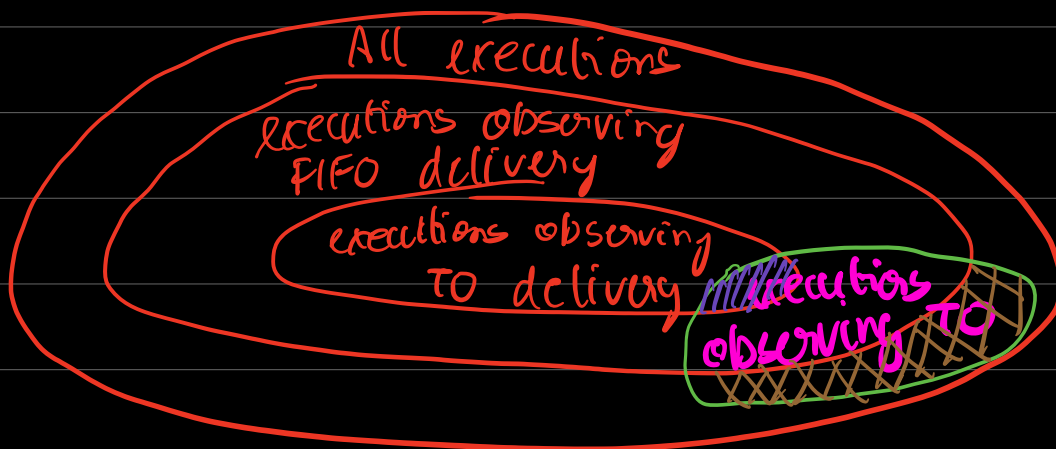
violation of FIFO  
delivery



violation of  
causal delivery



violation of TO  
delivery






Executions with TO + causal delivery



Executions with TO BUT NOT FIFO

An execution in the brown zone :



$P_2$  &  $P_3$  got  $m_2$  then  $m_1$ .  
However, the order of messages in both  $P_2$  &  $P_3$  was the SAME! Hence, it observes to delivery.

Implementing causal broadcast

Unicast messages : 1 sender 1 receiver  
(point-to-point)

Multicast messages : 1 sender many receivers


Broadcast : special case of multicast msgs  
1 sender everyone receives  
↓  
including sender itself!

## Vector clocks algorithm [with a twist]

don't count msg  
receives as  
events!


Don't count message receives as events!

- Every process keeps a VC, init at 0
- When a process sends a message, it increments its own position in its VC, and it's going to include updated VC with the message
- When a process **delivers** a message, it updates its VC to the pointwise max of its local VC & the received VC on the message



Causal delivery Property of executions that we care about today

Causal broadcast Algo that gives causal delivery in a setting where all messages are broadcast messages



Brown event is the problematic message!  
 $m_2$  is getting delivered before  $m_1$

How to use the vc to determine which message can be delivered now & which message should be reserved & delivered later?

Causal Broadcast

Define a deliverability condition that tells us if a received message is or is not OK to deliver



This deliverability condition will use the vector clock on the message.


## Deliverability:

A message  $m$  is deliverable at process  $P$  if, for  $k: 1 \dots n$  where  $VC$ 's have  $n$  entries

$$VC(m)[k] = VC(p)[k] + 1 \quad k \text{ is the sender}$$

$$VC(m)[k] \leq VC(p)[k] \quad \text{otherwise}$$





Queue size? Not too large

Snapshots of distributed systems

Global state of the system: How?

- Using globally synced clock can't use

Property that we want:

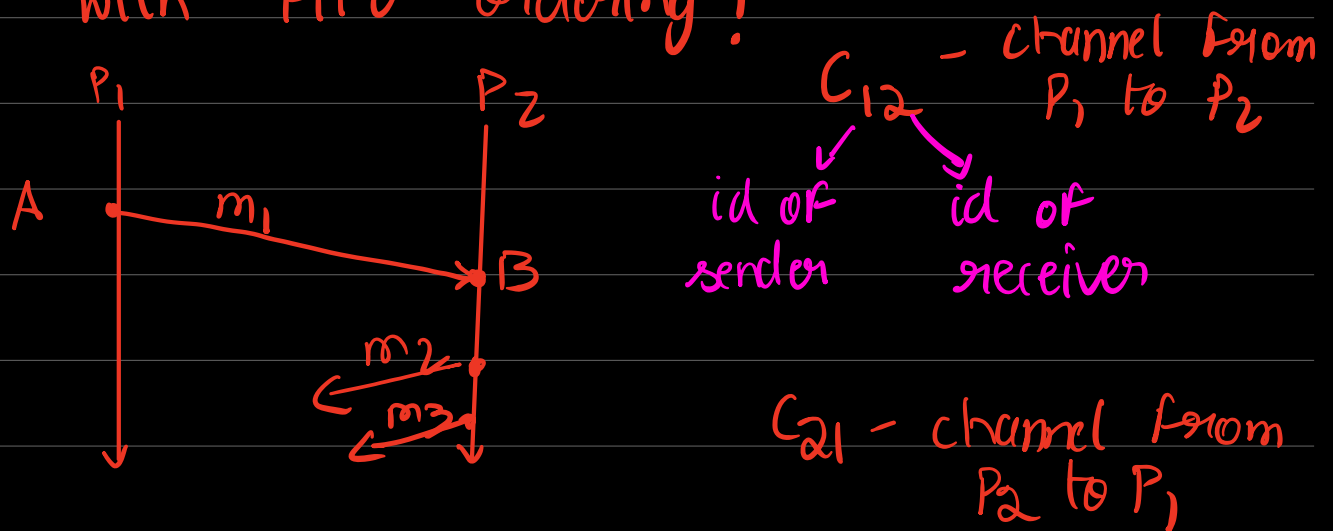
If we have events  $A$  &  $B$  where  $A \rightarrow B$  &  $B$  is in the snapshot,  $A$  is in the snapshot too!

Chandy - Lamport Algorithm

1985 algorithm

Channel: Connection from one process to another

with FIFO ordering!



Inflight msg ( $m_2$ ) are in the channel

$$C_{21} = [m_2, m_3]$$