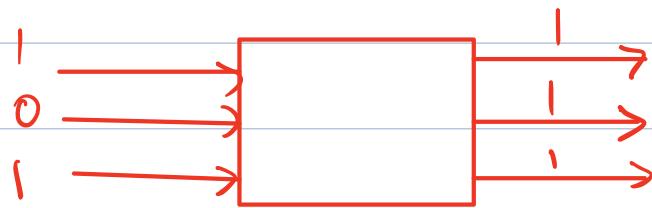


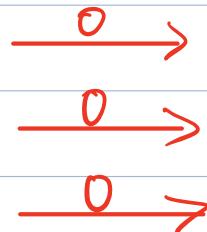
Agenda:

- definition of consensus
- FLP result
- Paxos: the easy parts
- Paxos: the interesting/hard parts

Consensus



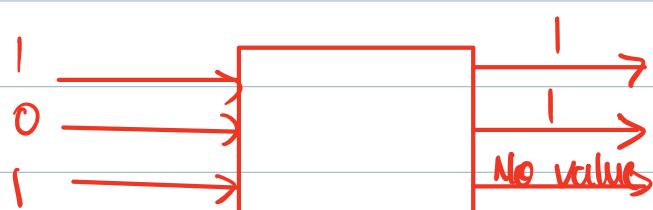
OR



Can be usually done with 'majority'

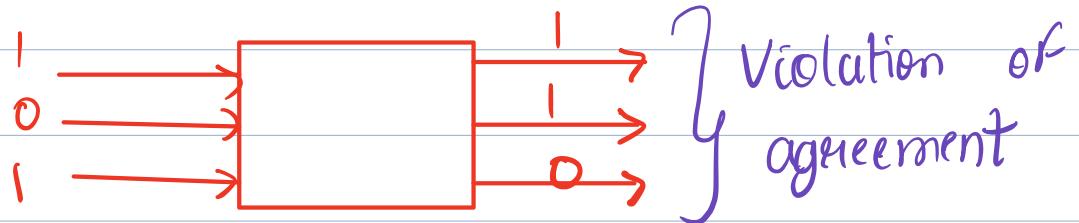
Properties that consensus algos TRY to satisfy:

Termination: Each correct process eventually decides on a value.



} Violation of
termination.

Agreement: All correct processes decide on the same value.

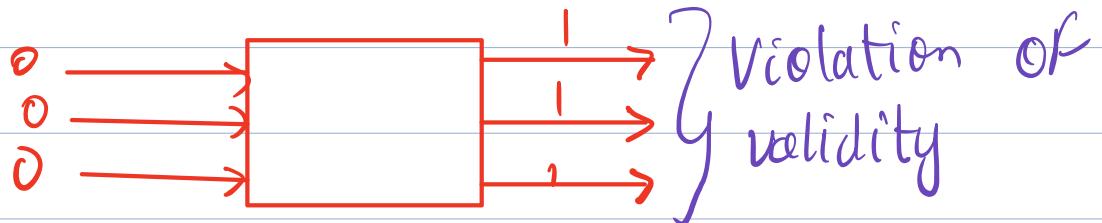


If only Termination & Agreement are the reqd conditions, how can we implement consensus easily?

Just have everyone output a constant value,
i.e. Ignore the input.

Validity(Integrity; non triviality):

The agreed upon value must be one of the proposed values



No consensus algo satisfies all 3 conditions in a sync network mode & a crash fault model. Why?
Impossible. Proved in 1983 by

Nancy Lynch
Michael Paterson

FLP Result!

In some cases, we can compromise on TERMINATION

Paxos Compromises on TERMINATION

If a process doesn't terminate, does it satisfy agreement?
YES!

PAXOS

Consensus Algorithm by Leslie Lamport (1998)
2001 - Paxos made simpl

Roles:-

Proposer : Propose values

Accepter : Contribute to choosing from among the proposed value

Learner - Learns the agreed upon value

One process could take on multiple roles

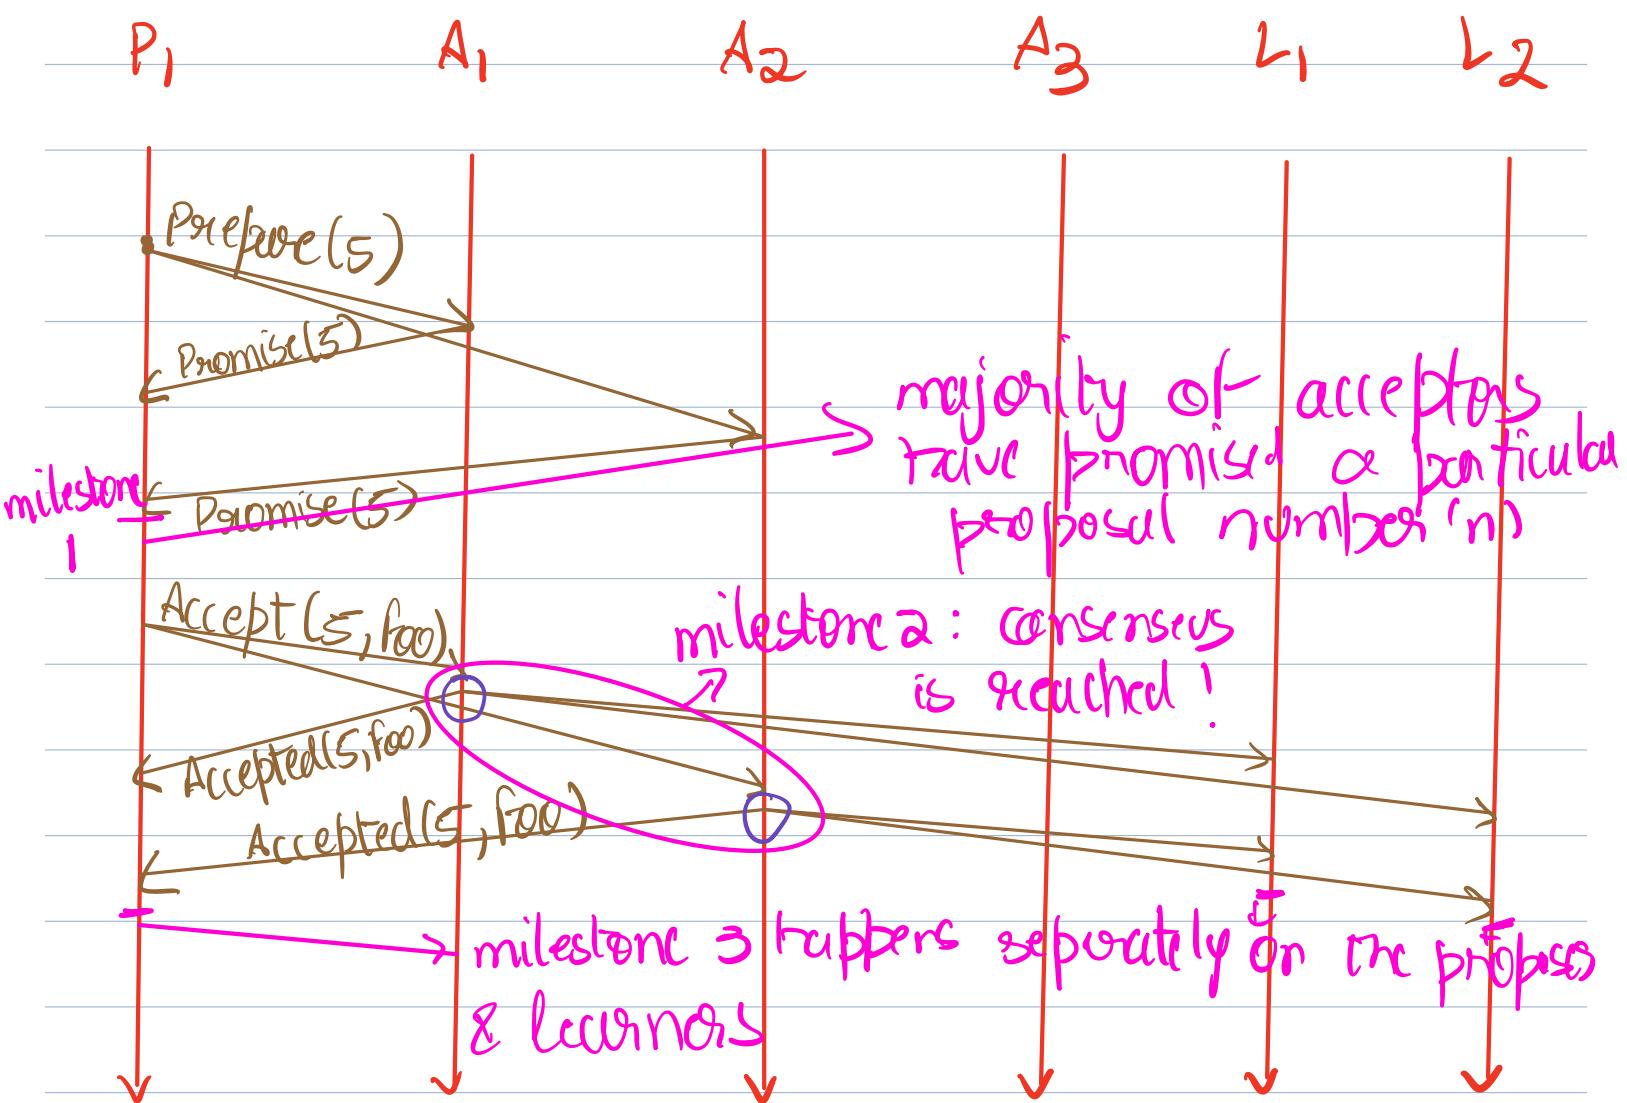
Paxos Node: any node that plays any role

Prerequisites

Paxos nodes must:

- Persist Data
- Know how many nodes is a majority of acceptors

Demo



- Majority Acceptors: 2

- P_i sends prepare msg to atleast majority of acceptors with a proposal number (n)
 - ' n ' has to be unique & higher than any number used by the current proposer but used before.
 - Store proposal number so as to not reuse them. Also the reason why nodes need to persist data

- When acceptor gets a proposal number, it will check to see if it previously promised to ignore this proposal number.
 - If yes, ignore

- If no : It now promises to ignore any proposal with number lower than ' n '. And reply to the proposer with a \star Promise(n) message
(Addressed later) ↓

I will ignore and greatest with proposal

number lower than ' n '.

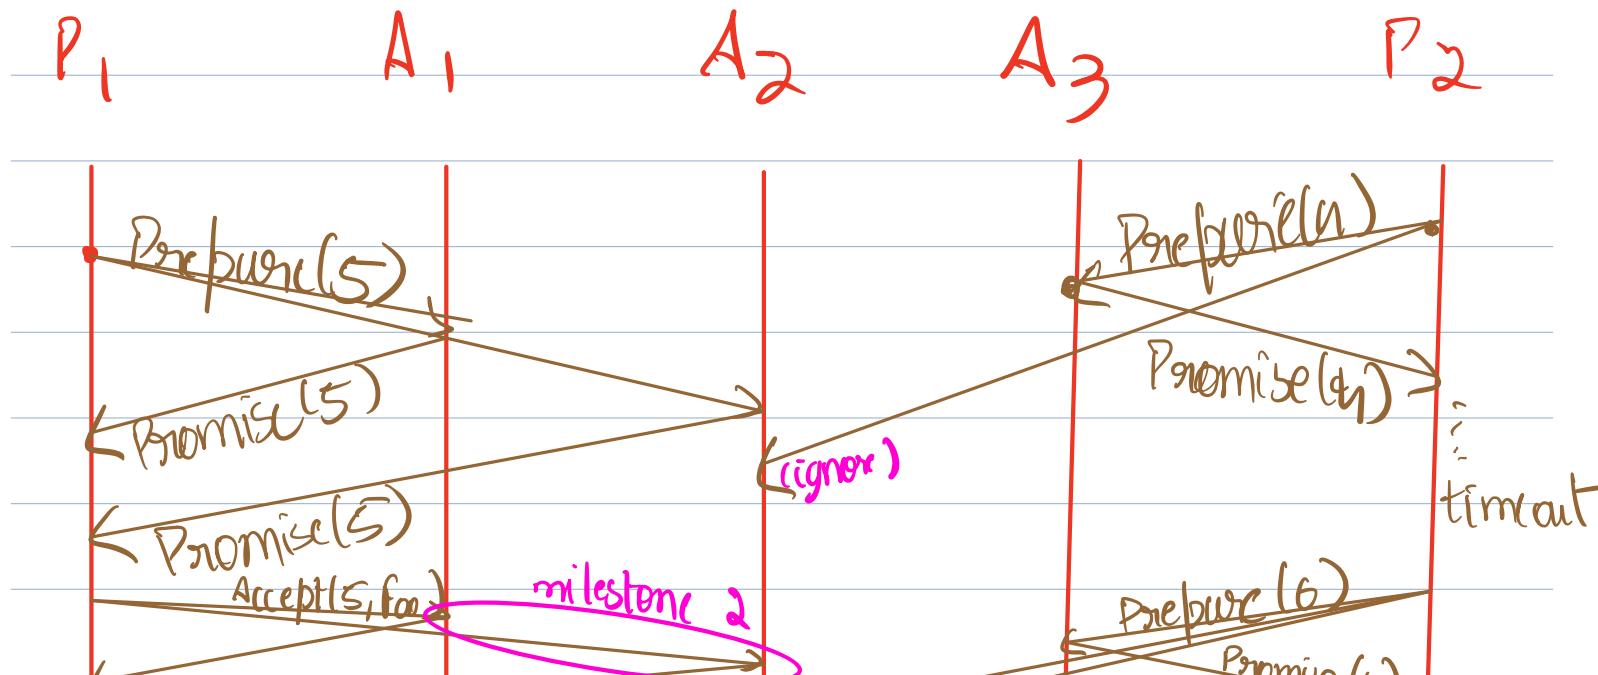
Paxos Phase 2 :-

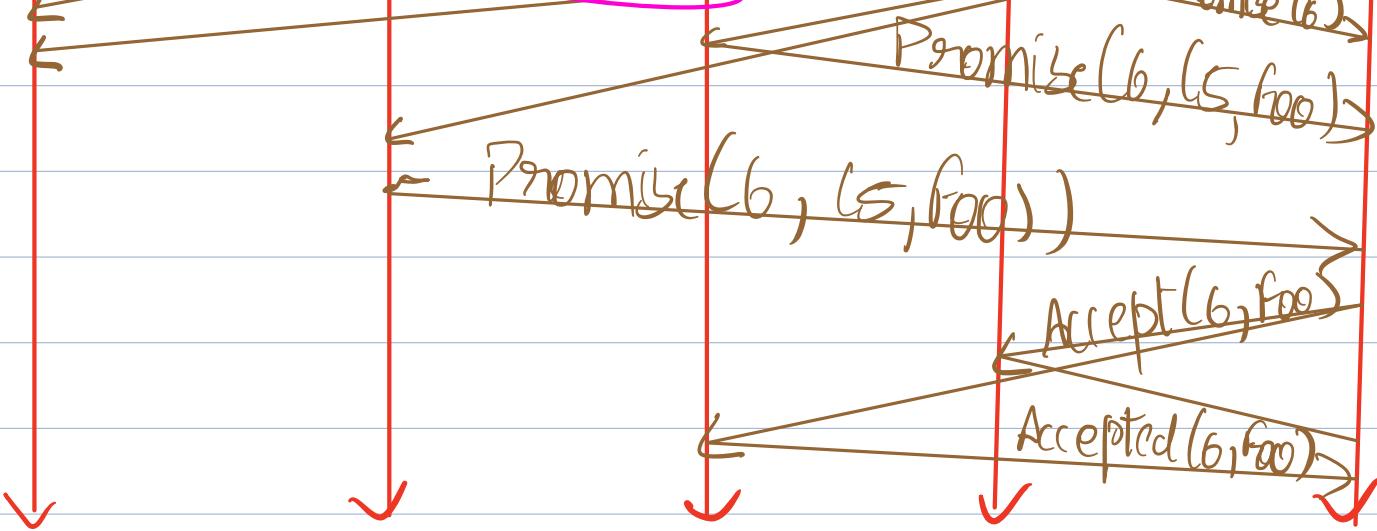
- Proposer has received $\text{Promise}(n)$ from a majority of acceptors for some ' n '.
- Send $\text{Accept}(n, \text{val})$ message to at least a majority of acceptors where
 - n - proposal number that was promised
 - val - Actual value it wants to propose
** Addressed later
- When acceptor gets the accept message : i.e $\text{Accept}(n, \text{val})$; it will check if it previously promised to ignore requests with this proposal number.
 - If yes; ignore message !
 - If no : Reply with $\text{Accepted}(n, \text{val})$, & also sends that message to all the Learners
(not just majority)

- The point at which majority of acceptors send out an Accepted (n, val) message, is when consensus is reached (milestone 2)
- When each individual process receives Accepted (n, val) message from majority of acceptors, they are now aware that consensus has been reached.

Addressing Asterisk (*) points:

- Let's assume we have multiple proposers





First Asterisk (*)

Acceptor :

On receiving a Prepar(n) message:

"Did I previously promise to ignore requests with proposal number?

- If yes: Ignore it
- If not: it asks

"Have I previously accepted anything?"

If yes: it replies with

KEY difference!

Promise(n , n_{prev} , val_{prev})

n_{prev} - highest previously accepted proposal number

val_{prev} : previously accepted value

If not: Reply with Promise(n)

Second asterisk (**)

(or PromiseIn(n_{prev} , val_{prev}))

- Proposer has received Promise(n) from a majority of acceptors for some ' n '.
- Send Accept(n , val) message to at least a majority of acceptors where

n - proposal number that was promised

val - Actual value it wants to propose
** (Addressed later)

val is chosen as follows:

The val prev that went with highest nprev

Key difference:

OR

If there are none, do
whatever it wants.

Consensus properties:

termination - when will Paxos
not terminate?
Agreement
Validity