

Agenda

- Online vs offline systems
raw vs derived data
- Intro to MapReduce
- Example: inverted index from forward index

Online Systems or Services

KV stores, web servers, databases, caches

wait for requests from clients, & try to handle them quickly.

Low Latency & Availability

Offline Systems or Batch processing Systems

MapReduce

Process LOTS of data

High throughput often priority

Streaming Systems

Hybrid of offline & online

Jamie Brandon

Different Representations of data

Raw data - Authoritative version

The format new data goes into

Derived data - Result of taking existing data & processing it somehow

MapReduce is a tool for computing derived data

Inverted Index

Inversion of forward index.

Forward Index

Inverted Index

DOC	word	word	Doc
DOC1	the, quick, brown, fox	the	Doc1, Doc2
Doc2	the, dog, grows	quick	Doc1
		dog	Doc2.

Converting Forward Index to Inverted Index.

for each document D:

for each word w in D :

emit $\langle w, D \rangle$

then,

combine DS in a list for
each unique word w

why can't we use above pseudocode in real life?

At scale, this become a distributed systems
problem, & needs to deal with things like
partitioning, availability, etc.

MapReduce is a framework to enable solving
these problems, by hiding the complexity of
distributed systems.

for each document D :

This can be done in parallel!

for each word w in D :

emit $\langle w, D \rangle$

This can be done
locally on a
machine for a
doc

m_1

$\langle \text{Doc1}, [\text{the}, \text{quick} \dots] \rangle$

$\langle \text{the}, \text{Doc1} \rangle$

$\langle \text{quick}, \text{Doc1} \rangle$

m_2

$\langle \text{Doc2}, [\text{the}, \text{dog} \dots] \rangle$

Intermediate
KV pairs

R_1

$\langle \text{the}, \text{Doc1} \rangle$

$\langle \text{the}, \text{Doc2} \rangle$

m_3

$\langle \text{the, Doc 2} \rangle$

$\langle \text{dog, Doc 2} \rangle$

$\langle \text{Doc 3, [i, love, my...]} \rangle$

$\langle \text{i, Doc 3} \rangle$

$\langle \text{love, Doc 3} \rangle$

$\langle \text{dog, Doc 2} \rangle$

$\langle \text{dog, Doc 3} \rangle$

R_2

$\langle \text{quick, Doc 1} \rangle$

$\langle \text{love, Doc 3} \rangle$

\Downarrow
 $\langle \text{the, [Doc 1, Doc 2]} \rangle$

How would you aggregate intermediate KV pairs

Ensure all same keys end up on same machine.


Implementing using HASHING

$\text{hash}(\text{key}) \% 2$

\downarrow
No. of 'reduce' machines

How does 'Reduce' worker recall the data?

In general, every machine needs to read data from all 'map' nodes.



This communication pattern is called shuffle

Examples of tasks for MapReduce

- inverted index
- grep
- distributed sorting
- distributed word count

Combiners

Run the 'reduce' logic locally on the map workers & store intermediate data.